

Boqing Gong, Ph.D.

CONTACT INFORMATION	✉ bgong@bu.edu	🌐 https://boqinggong.github.io
	☎ +1(213)810-7118	CDS 822, 665 Commonwealth Ave, Boston, MA 02215
APPOINTMENTS	Assistant Professor in Computer Science at Boston University (BU)	07/2024 –
	<i>Research, teaching, services, NSF CAREER Award</i>	
	Secondary appointment in System Engineering	08/2025 – 08/2028
	Junior Faculty Fellow, Faculty of Computing & Data Science	08/2025 – 08/2028
	Junior Faculty Fellow, Rafik B. Hariri Institute	05/2026 – 05/2029
	Staff Research Scientist at Google DeepMind (part-time since 09/2024)	03/2019 –
	<i>Research and productionization of computer vision and deep learning</i>	
	Adjunct Lecturer at Brown University (part-time)	01/2021 – 04/2021
	<i>Taught Deep Learning, overall rating 4.33 vs. department mean 3.94</i>	
	Principal Investigator (PI) at International Computer Science Institute, Berkeley	01/2018 – 12/2020
	<i>Research on computer vision, NSF Award #1835539</i>	
	Principal Researcher at Tencent America	01/2018 – 03/2019
EDUCATION	<i>Research on computer vision and reinforcement learning</i>	
	Assistant Professor at University of Central Florida (UCF)	08/2015 – 12/2017
	Graduate Faculty Member	08/2015 –
	<i>Research on computer vision and machine learning, PI for NSF Awards 1566511 and 1741431</i>	
	<i>Supervision of six Ph.D. students and numerous master and undergraduate students</i>	
	<i>Teaching: advanced topics in computer vision, robot vision, etc.</i>	
	<i>Services: departmental committees, NSF panels, and conference organization committees</i>	
	Summer Research Assistant at NEC Laboratories America (internship)	06/2013 – 08/2013
	<i>Feature engineering for large-scale, fine-grained object recognition</i>	
	Research Assistant at Microsoft Research Asia (internship)	01/2008 – 06/2008
	<i>Feature engineering for content-based image retrieval</i>	
	Ph.D. in Computer Science, University of Southern California, USA	08/2011 – 08/2015
EDUCATION	Thesis: Kernel Methods for Unsupervised Domain Adaptation	
	Thesis Committee: Fei Sha (supervisor), Gaurav Sukhatme, and Shrikanth Narayanan	
	Visiting Ph.D. Student in Computer Science, University of Texas at Austin, USA	Summer, 2013
	Host: Prof. Kristen Grauman	
	M.Phil. in Information Engineering, The Chinese University of Hong Kong	08/2008 – 07/2010
	Thesis: 3D Object Retrieval and Recognition	
	Thesis Committee: Xiaoou Tang (supervisor), Jianzhuang Liu, and Xiaogang Wang	
	University of Science and Technology of China	09/2004 – 07/2008
	B.E. in Electronic Engineering and Information Science	

INVITED TALKS

<i>BabyVFM-V2: Democratizing Research on Pretraining Vision Foundation Models</i>	
Indiana University Bloomington, Computer Science Department Colloquium	10/16/2026
Sony Japan	09/03/2026
Seoul National University, Host: Bohyung Han	08/18/2026
CCN 2026 Satellite Event on Beyond Curated Datasets	08/02/2026
CVPR 2026 Workshop on Open-World Vision	06/04/2026
University of Illinois Urbana-Champaign, Host: Shenlong Wang	05/01/2026
<i>BabyVFM: Data-Efficient Pretraining of Vision LLMs Inspired by Infant Learning</i>	
Egocentric Video Workshop at Stanford University, Host: Michael C. Frank	02/18/2026
Massachusetts Institute of Technology, Host: Kaiming He	12/19/2025
Northeastern University, Host: Huaizu Jiang	12/17/2025
University of Maryland, College Park, Host: Ruohan Gao	12/16/2025
Texas A&M University, Host: Tianbao Yang	12/11/2025
University of Texas at Austin, Host: Qixing Huang	12/10/2025
Johns Hopkins University, Host: Zongwei Zhou and Alan L. Yuille	10/31/2025
BU MET Computer Science Research Seminar Series, Host: Reza Rawassizadeh	10/10/2025
<i>Democratizing Research on the Pretraining of Vision Large Language Models</i>	
CVPR 2025 Workshop on Computer Vision in the Wild	06/11/2025
University of Texas at Dallas, Host: Yunhui Guo	05/01/2025
<i>From Domain Adaptation to Generalist Vision Systems: A Quest for Out-of-Domain Generalization</i>	
ACCV Workshop on LAMM	12/09/2024
Georgia Institute of Technology, Host: Bo Dai	11/08/2024
University of California, Merced, Host: Ming-Hsuan Yang	09/16/2024
ByteDance Foudamental Research, Host: Lu Jiang	08/07/2024
Microsoft Research at Redmond, Host: Reuben Tan	08/01/2024
Northwestern University, Host: Ulas Bagci	07/25/2024
Korea Advanced Institute of Science & Technology, Host: Sung Ju Hwang	04/19/2024
<i>Efficient Video Models</i>	
Guest lecture at Georgia Tech, hosted by Judy Hoffman	10/20/2022
Guest lecture at University of California Merced, hosted by Ming-Hsuan Yang	11/22/2021
<i>When Vision Transformers Outperform ResNets</i>	
ACM MM 2021 Workshop: Adversarial Learning for Multimedia	10/20/2021
CVPR 2021 Workshop: Learning from Limited and Imperfect Data	06/20/2021
<i>Data-Efficient Learning of Vision Transformers</i>	
Guest lecture at UNC Chapel Hill, host by Zhengwu Zhang	03/05/2024
Guest lecture at Georgia Tech, hosted by Judy Hoffman	10/20/2022
Computer Science Colloquium Talk, University of Illinois Urbana-Champaign	03/23/2022
ACML 2021 Workshop: Weakly Supervised Learning	11/16/2021
Rutgers University Efficient AI Seminar	08/10/2021
<i>Long-Tailed Visual Recognition in the Wild: A Tutorial</i>	
IJCAI 2021 Workshop: Long-Tailed Distribution Learning	08/21/2021
CVPR 2021 Tutorial: Data- and Label-Efficient Learning in An Imperfect World	06/19/2021
<i>Long-Tailed Visual Recognition in the Wild</i>	

ECCV 2020 Workshop on Imbalance Problems in Computer Vision	08/28/2020
<i>Towards Visual Recognition in the Wild: Long-Tailed Sources and Open Compound Targets</i>	
Center for Language and Speech Processing at Johns Hopkins University	12/04/2020
Computer Vision Group at University of Bristol	10/13/2020
Visual Informatics Group at University of Texas, Austin	09/11/2020
IEEE CVPR 2020 Workshops on Learning from Imperfect Data	06/14/2020
IEEE CVPR 2020 Workshop on Adversarial Machine Learning in Computer Vision	06/19/2020
<i>Long-Tailed Visual Recognition is A Domain Adaptation Problem</i>	
WACV Workshop on Vision Applications and Solutions to Biased or Scarce Data	03/05/2020
Google Research Conference	02/26/2020
IEEE CVPR 2020 Area Chair Workshop	01/24/2020
<i>NATTACK by Learning the Distributions of Adversarial Examples</i>	
Vision and Learning Seminar (Valse)	09/25/2019
IEEE CVPR Workshop on GigaVision	06/17/2019
Waymo Inc.	06/05/2019
Department of Computer Science, UC Davis	05/17/2019
<i>Curriculum Domain Adaptation</i>	
IEEE BIGDATA Workshop on Big Data Transfer Learning	12/10/2018
<i>The Multiple Shades of Dropout for Discriminative and Generative Deep Neural Networks</i>	
INFORMS Special Session on Stochastic Optimization Methods and Approximation Theory in Machine Learning	11/04/2018
<i>Domain Adaptation and Transfer: All You Need to Use Simulation “for Real”</i>	
ECCV Workshop on Visual Learning and Embodied Agents in Simulation Environments	09/09/2018
<i>Learning and Adapting from the Web for Visual Recognition</i>	
ECCV Workshop on Compact and Efficient Feature Representation and Learning in CV	09/09/2018
IEEE CVPR Workshop on Visual Understanding by Learning from Web Data	06/18/2018
<i>Domain Adaptation for Robust Visual Recognition and Semantic Segmentation</i>	
The Computer Vision Group at University of California, Merced	10/26/2020
Shenzhen Institutes of Advanced Technology, Chinese Academy of Sciences	04/02/2018
Shenzhen University	03/26/2018
International Computer Science Institute, UC Berkeley	12/08/2017
Beijing University of Posts and Telecommunications	11/06/2017
Department of Electrical Engineering, UC Santa Cruz	10/25/2017
Google Research at Mountain View	09/20/2017
Department of Media Analytics, NEC Laboratories America	05/08/2017
NVIDIA Research	06/08/2017
<i>Sequential Determinantal Point Processes for Supervised Video Summarization</i>	
Department of Computer Science, Stanford University	03/20/2017
Adobe Systems Inc.	03/30/2017
Facebook Inc.	06/07/2017
University of California at Berkeley	08/24/2017
<i>Domain Adaptation for Human Activity Recognition and Summarization</i>	
Army Research Office / Information Science Institute Workshop on Multi-Modal Data Analysis for	

	Human Activity Detection and Understanding	09/13/2016
	<i>Query-Focused Extractive Video Summarization</i>	
	Electrical Engineering and Computer Sciences, Univ. California at Berkeley	09/21/2017
	Department of Computer Science, University of California at Irvine	07/08/2016
	Snapchat Inc.	08/18/2016
	<i>Kernel Methods for Unsupervised Domain Adaptation</i>	
	Information Science Institute, University of Southern California	12/11/2015
	Department of Computer Science, Tulane University	04/23/2015
	Department of Machine Learning, NEC Laboratories America	04/09/2015
	Department of EECS, University of Central Florida	04/07/2015
	School of Computing, Informatics, and Decision Systems Engineering, ASU	04/02/2015
	IBM T.J. Watson Research Center (colloquium)	01/15/2015
	ECCV Workshop on TASK-CV	09/12/2014
	<i>Reshaping Datasets for Unsupervised Domain Adaptation</i>	
	IEEE ICDM Workshop on Practical Transfer Learning	11/14/2015
	<i>Sequential Determinantal Point Process: Modeling Video's Diverse and Sequential Properties</i>	
	Department of EECS, University of Central Florida	07/08/2015
	<i>Discriminative Kernel Learning for Unsupervised Domain Adaptation</i>	
	Machine Learning and Instrument Autonomy Group, JPL, NASA	01/09/2014
TUTORIALS	<i>BabyVLM: Democratizing Research on the Pretraining of Vision Large Language Models</i>	
	Half day, at IEEE International Conference on Development and Learning	09/15/2026
	<i>Transformers for Multimodal Signal Processing and Decision Making</i>	
	A 10-hour course at Int'l Conference on Acoustics, Speech, and Signal Processing	05/26/2022
	A short version at Google-UNLV ExploreCSR Workshop	05/25/2023
	<i>How to Train Your (Vision) Transformer</i>	
	AAAI 2023 Tutorial on Everything You Need to Know about Transformers	02/07/2023
	<i>Sequential Determinantal Point Processes: Models, Algorithms, and Applications</i>	
	CVPR Tutorial on Recent Advances in Visual Data Summarization	06/16/2019
ACADEMIC & PROFESSIONAL SERVICES	National Science Foundation panelist	
	2026 (1 panel), 2025 (1), 2024 (3), 2023 (2), 2021 (3), 2020 (2), 2019 (1), 2017 (1), 2016 (3)	
	Associate editor	
	Springer International Journal on Computer Vision (IJCV)	08/2025 –
	IEEE Transactions on Pattern Analysis and Machine Intelligence (T-PAMI)	06/2024 –
	Transactions on Machine Learning Research (TMLR)	12/2024 –
	Program co-chair	
	IEEE Winter Conference on Applications of Computer Vision (WACV)	2023, 2028
	NeurIPS Workshop on BabyVLM	2026
	CVPR Workshop on Multi-Modal Learning from Videos	2019
	Google Pre-CVPR 2020 Workshop	2020
	Google Mobile Vision Workshop	2020
	Workshop co-chair	

	IEEE International Conference on Computer Vision (ICCV)	2025
	Tutorial co-chair	
	IEEE Conference on Computer Vision and Pattern Recognition (CVPR)	2022, 2026
	Lead guest editor	
	Machine Vision and Applications Special Issue on WACV 2023 Award Finalists	2022 – 2023
	Senior area chair	
	International Conference on Machine Learning (ICML)	2026
	IEEE Conference on Computer Vision and Pattern Recognition (CVPR)	2026, 2027
	International Conference on Learning Representations (ICLR)	2025 – 2027
	IEEE International Conference on Computer Vision (ICCV)	2025
	AAAI Conference on AI	2020, 2021, 2024, 2026
	International Joint Conference on Artificial Intelligence (IJCAI)	2021
	Area chair / Meta reviewer	
	International Conference on Learning Representations (ICLR)	2021 – 2024
	IEEE Conference on Computer Vision and Pattern Recognition (CVPR)	2020, 2022 – 2024
	European Conference on Computer Vision (ECCV)	2020, 2022
	Neural Information Processing Systems (NeurIPS)	2019 – 2024
	IEEE International Conference on Computer Vision (ICCV)	2019, 2021, 2023, 2025
	IEEE Winter Conference on Applications of Computer Vision (WACV)	2018 – 2020, 2024
	International Conference on Machine Learning (ICML)	2019 – 2025
	International Conference on Artificial Intelligence and Statistics (AISTATS)	2019
	Mentor: WACV 2018 Ph.D. Forum, ICCV 2021 Doctoral Consortium	
	Chair: IEEE Signal Processing's Continuing Education for Industry Subcommittee	2023 – 2024
	Publicity co-chair	
	IEEE Winter Conference on Applications of Computer Vision (WACV)	2024
	IEEE Conference on Computer Vision and Pattern Recognition (CVPR)	2023, 2024
	IEEE International Conference on Computer Vision (ICCV)	2023
	Reviewer: NeurIPS, ICML, ICLR, CVPR, ICCV, ECCV, AISTATS, JMLR, IJCV, T-PAMI...	
	Workshop panelist	
	CVPR: Computer Vision for Children	06/04/2026
	CVPR: Computer Vision in the Wild	06/19/2023
	ICLR: What do we need for successful domain generalization?	05/04/2023
	CVPR: Learning from Limited or Imperfect Data	06/20/2021
	ECCV: Imbalance Problems in Computer Vision	08/28/2020
	CVPR: Adversarial Machine Learning in Computer Vision	06/19/2020
	ECCV: Visual Learning and Embodied Agents in Simulation Environments	09/09/2019
DEPARTMENTAL SERVICES	CS Faculty Search Committee (BU)	2024 – 2025
	Graduate Awards Committee (BU)	2025 –
	CS Faculty Search Committee (UCF)	2017 – 2018
	Awards Committee, College of Engineering and Computer Science (UCF)	2017 – 2018
	CRCV Research Associate Search Committee (UCF)	2016
	Nielsen Fellowship Search Committee (UCF)	2016

SELECTED
PUBLICATIONS

Statistics as of 09/09/2026 according to Google Scholar:
Citations: 29,600+ h-index: 63 i10-index: 109 citations to (CVPR'12) [C4]: 3,200+

The complete publication list is available [here](#).

= Equal contribution

INVITED BOOK CHAPTERS

- [B2] **B. Gong**, K. Grauman, and F. Sha. “Geodesic Flow Kernel and Landmarks: Kernel Methods for Unsupervised Domain Adaptation.” In *Domain Adaptation for Computer Vision Applications*, Springer Publishing, 2017.
- [B1] C. Gan*, T. Yang, and **B. Gong**. “A Multi-Source Domain Generalization Approach to Visual Attribute Detection.” In *Domain Adaptation for Computer Vision Applications*, Springer Publishing, 2017.

JOURNAL PUBLICATIONS

- [J8] L. Yuan, N. B. Gundavarapu, L. Zhao, H. Zhou, Y. Cui, L. Jiang, X. Yang, M. Jia, T. Weyand, L. Friedman, M. Sirotenko, H. Wang, F. Schroff, H. Adam, M.-H. Yang, T. Liu, and **B. Gong**. “Videoglue: Video General Understanding Evaluation of Foundation Models.” *Transactions on Machine Learning Research (TMLR)*, 2024.
- [J7] Z. Liu, Z. Miao, X. Zhan, J. Wang, **B. Gong**, S. Yu. “Open Long-Tailed Recognition in A Dynamic World.” *IEEE Transactions on Pattern Analysis and Machine Intelligence (T-PAMI)*, 2022.
- [J6] Y.-C. Su, S. Changpinyo, S. Thoppay, C.-J. Hsieh, L. Shapira, R. Soricut, H. Adam, M. Brown, M.-H. Yang, and **B. Gong**. “2.5D Visual Relationship Detection.” *Computer Vision and Image Understanding (CVIU)*, 2022.
- [J5] S. Changpinyo, W.-L. Chao, **B. Gong**, and F. Sha. “Classifier and Exemplar Synthesis for Zero-Shot Learning.” *International Journal of Computer Vision (IJCV)*, 2019.
- [J4] Y. Zhang*, P. David, F. Hassan, and **B. Gong**. “A Curriculum Domain Adaptation Approach to the Semantic Segmentation of Urban Scenes.” *IEEE Transactions on Pattern Analysis and Machine Intelligence (T-PAMI)*, 2019.
- [J3] A. Mazaheri*, **B. Gong**, and M. Shah. “Learning a Multi-Concept Video Retrieval Model with Multiple Latent Variables.” *ACM Transactions on Multimedia Computing, Communications, and Applications (TOMM)*, Vol. 14, Issue 2, May 2018.
- [J2] **B. Gong**, K. Grauman, and F. Sha. “Learning Kernels for Unsupervised Domain Adaptation with Applications to Visual Object Recognition.” *International Journal of Computer Vision (IJCV)*, Vol. 109, Issue 1-2, pp. 3-27, August 2014. [[Link](#)]
- [J1] **B. Gong**, J. Liu, X. Wang, and X. Tang. “Learning Semantic Signatures for 3D Object Retrieval.” *IEEE Transactions on Multimedia (T-MM)*, Vol. 5, Issue 2, pp. 369-377, February 2013.

PEER-REVIEWED CONFERENCE PUBLICATIONS

- [C97] M. Whitton, Z. Wang, P. Liu, Q. T. Truong, S. Wang, M. Yadamreddy, O. Ozel, V. Jayanti, S. Sekhon, H. S. Tadesse, L. Miao, J. Wang, J. Lu, C. Yu, and **B. Gong**. “Making Sense of Touch from the Child’s View for Contrastive Learning.” *Proceedings of the IEEE International Conference on Development and Learning (ICDL)*, Japan, September 2026.
- [C96] Y. Qing, C. Mao, and **B. Gong**. “StochasT: Learning with Stochastic Turn Depth for Visual Instruction Tuning.” *Proceedings of the European Conference on Computer Vision (ECCV)*, Sweden, September 2026.

- [C95] Y. Tan, J. Huang, J. Huang, H. Li, and **B. Gong**. “Natural Language Camera Movement Understanding.” *Proceedings of the European Conference on Computer Vision (ECCV)*, Sweden, September 2026.
- [C94] Y. Tan, Y. Qing, and **B. Gong**. “The LLM Bottleneck: Why Open-Source Vision LLMs Struggle with Hierarchical Visual Recognition.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Denver, June 2026.
- [C93] S. Wang⁺, W. Wang⁺, Z. Wang⁺, M. Whitton⁺, M. Wakeham, A. Chandra, J. Huang, P. Zhu, H. Chen, D. Li, J. Li, S. Li, A. Zagula, A. Zhao, A. Zhu, S. Nakamura, Y. Yamamoto, J. J. Yokono, A. Mueller, B. A. Plummer, K. Saenko, V. Saligrama, and **B. Gong**. “BabyVLM-V2: Toward Developmentally Grounded Pretraining and Benchmarking of Vision Foundation Models.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Denver, June 2026.
- [C92] S. Wang^{*}, A. Chandra, A. Liu, V. Saligrama, **B. Gong**. “BabyVLM: Data-Efficient Pretraining of VLMs Inspired by Infant Learning.” *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, Honolulu, HI, 2025.
- [C91] Z. Zhang^{*}, M. Dou, L. Peng, H. Pan, U. Bagci, **B. Gong**. “VideoAds for Fast-Paced Video Understanding.” *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, Honolulu, HI, 2025.
- [C90] W. Wang^{*}, R. Tan, P. Zhu, J. Yang, Z. Yang, L. Wang, A. Kolobov, J. Gao⁺, **B. Gong**⁺. “SITE: towards Spatial Intelligence Thorough Evaluation.” *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, Honolulu, HI, 2025.
- [C89] L. Li^{*}, K. Fan, **B. Gong**⁺, X. Yue⁺. “HypDAE: Hyperbolic Diffusion Autoencoders for Hierarchical Few-shot Image Generation.” *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, Honolulu, HI, 2025.
- [C88] Z. Zou^{*}, **B. Gong**, L. Wang. “Attention to Neural Plagiarism: Diffusion models Can Plagiarize Your Copyrighted Images!” *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, Honolulu, HI, 2025.
- [C87] L. Zhao, S. Woo, Z. Wan, Y. Li, H. Zhang, **B. Gong**, H. Adam, X. Jia, T. Liu. “ ϵ -VAE: Denoising as Visual Decoding.” *Proceedings of the International Conference on Machine Learning (ICML)*, Vancouver, Canada, 2025.
- [C86] L. Chen^{*}, H. Hu, M. Zhang, Y. Chen, Z. Wang, Y. Li, P. Shyam, T. Zhou, H. Huang, M.-H. Yang, **B. Gong**. “OmnixR: Evaluating Omni-modality Language Models on Reasoning Across Modalities.” *Proceedings of the International Conference on Learning Representations (ICLR)*, Singapore, April 2025.
- [C85] Y. Ban, R. Wang, T. Zhou, **B. Gong**, C.-J. Hsieh, M. Cheng. “The Crystal Ball Hypothesis in diffusion models: Anticipating object positions from initial noise.” *Proceedings of the International Conference on Learning Representations (ICLR)*, Singapore, April 2025.
- [C84] N. B. Gundavarapu, L. Friedman, R. Goyal, C. Hegde, E. Agustsson, S. M. Waghmare, M. Sirotenko, M.-H. Yang, T. Weyand, **B. Gong**, L. Sigal. “Extending Video Masked Autoencoders to 128 Frames.” *Proceedings of the Neural Information Processing Systems (NeurIPS)*, Canada, December 2024.
- [C83] Y. Ban, R. Wang, T. Zhou, M. Cheng, **B. Gong**, C.-J. Hsieh. “Understanding the Impact of Negative Prompts: When and How Do They Take Effect?” *Proceedings of the European Conference on Computer Vision (ECCV)*, Italy, September 2024.
- [C82] L. Zhao⁺, N.B. Gundavarapu⁺, L. Yuan⁺, H. Zhou⁺, Sh. Yan, J.J. Sun, L. Friedman, R. Qian, T. Weyand, Y. Zhao, R. Hornung, F. Schroff, M.-H. Yang, D.A. Ross, H. Wang, H. Adam, M. Sirotenko⁺, T. Liu⁺, and **B. Gong**⁺. “VideoPrism: A Foundational Visual Encoder for Video Understanding.” *Proceedings of the International Conference on Machine Learning (ICML)*, Austria, July 2024.

- [C81] R. Wang, T. Liu, C.J. Hsieh, and **B. Gong**. “dPOD: On Discrete Prompt Optimization for Diffusion Models.” *Proceedings of the International Conference on Machine Learning (ICML)*, Austria, July 2024.
- [C80] Y. Zhao, L. Zhao, X. Zhou, J. Wu, C.-T. Chu, H. Miao, F. Schroff, H. Adam, T. Liu, **B. Gong**, P. Krhenbhl, and L. Yuan. “Distilling vision-language models on millions of videos.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Seattle, June 2024.
- [C79] H. Hu, K. CK Chan, Y.-C. Su, W. Chen, Y. Li, K. Sohn, Y. Zhao, X. Ben, **B. Gong**, W. Cohen, M.-W. Chang, and X. Jia. “Instruct-Imagen: Image generation with multi-modal instruction.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Seattle, June 2024.
- [C78] L. Yu, J. Lezama, N. B Gundavarapu, L. Versari, K. Sohn, D. Minnen, Y. Cheng, A. Gupta, X. Gu, A. G Hauptmann, **B. Gong**, M.-H. Yang, I. Essa, D. A Ross, and L. Jiang. “Language Model Beats Diffusion—Tokenizer is Key to Visual Generation.” *Proceedings of the International Conference on Learning Representations (ICLR)*, Austria, May 2024.
- [C77] Y. Xiong, L. Zhao, **B. Gong**, M.-H. Yang, F. Schroff, T. Liu, C.-J. Hsieh, and L. Yuan. “Structured Video-Language Modeling with Temporal Grouping and Spatial Grounding.” *Proceedings of the International Conference on Learning Representations (ICLR)*, Austria, May 2024.
- [C76] C. Liang, J. Yu, M.-H. Yang, M. Brown, Y. Cui, T. Zhao, **B. Gong**, and T. Zhou. “Module-Wise Adaptive Distillation for Multimodality Foundation Models.” *Proceedings of the Neural Information Processing Systems (NeurIPS)*, New Orleans, December 2023.
- [C75] M. Liu, M. Zhang, J. Liu, H. Dai, M.-H. Yang, S. Ji, Z. Feng, and **B. Gong**. “Video Timeline Modeling for News Story Understanding.” *Proceedings of the Neural Information Processing Systems Datasets and Benchmarks Track*, New Orleans, December 2023.
- [C74] L. Zhao, L. Yuan, **B. Gong**, Y. Cui, F. Schroff, M.-H. Yang, H. Adam, and T. Liu. “Unified Visual Relationship Detection with Vision and Language Models.” *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, Paris, October 2023.
- [C73] D. Wang, **B. Gong**, and L. Wang. “On Calibrating Semantic Segmentation Models: Analyses and An Algorithm.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Vancouver, June 2023.
- [C72] M. Liu, Y. Zhou, C.-R. Qi, **B. Gong**, H. Su, and D. Anguelov. “LESS: Label-Efficient Semantic Segmentation for LIDAR Point Clouds.” *Proceedings of the European Conference on Computer Vision (ECCV)*, Tel Aviv, October 2022.
- [C71] Z. Zou*, **B. Gong**, and L. Wang. “Anti-Neuron Watermarking: Protecting Personal Data Against Unauthorized Neural Networks.” *Proceedings of the European Conference on Computer Vision (ECCV)*, Tel Aviv, October 2022.
- [C70] R. Qian, Y. Li, L. Yuan, **B. Gong**, T. Liu, M. Brown, S. Belongie, M.-H. Yang, H. Adam, and Y. Cui. “On Temporal Granularity in Self-Supervised Video Representation Learning.” *Proceedings of the British Machine Vision Conference (BMVC)*, 2022.
- [C69] X. Chen, C.-J. Hsieh, and **B. Gong**. “When Vision Transformers Outperform ResNets without Pre-training or Strong Data Augmentations.” *Proceedings of the International Conference on Learning Representations (ICLR)*, Virtual, April 2022.
- [C68] J. Zhuang, **B. Gong**, L. Yuan, Y. Cui, H. Adam, N.C. Dvornek, S. Tatikonda, J.S. Duncan, and T. Liu. “Surrogate Gap Minimization Improves Sharpness-Aware Training.” *Proceedings of the International Conference on Learning Representations (ICLR)*, Virtual, April 2022.
- [C67] C.-H. Yao, **B. Gong**, H. Qi, Y. Cui, Y. Zhu, and M.-H. Yang. “Federated Multi-Target Domain Adaptation.” *Proceedings of the Winter Conference on Applications of Computer Vision (WACV)*, Waikoloa, HI, Jan 2022.

- [C66] L. Yuan, R. Qian, Y. Cui, **B. Gong**, F. Schroff, M-H. Yang, H. Adam, and T. Liu. “Contextualized Spatio-Temporal Contrastive Learning with Self-Supervision.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2022.
- [C65] A. Akula, S. Changpinyo, **B. Gong**, P. Sharma, S.-C. Zhu, and R. Soricut. “CrossVQA: Scalably Generating Benchmarks for Systematically Testing VQA Generalization.” *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, Virtual, November 2021.
- [C64] A. Iscen, A. Araujo, **B. Gong**, and C. Schmid. “Class-Balanced Distillation for Long-Tailed Visual Recognition.” *Proceedings of the British Machine Vision Conference (BMVC)*, Virtual, November 2021.
- [C63] H. Akbari, L. Yuan, R. Qian, W.-H. Chuang, S.F. Chang, Y. Cui, and **B. Gong**. “VATT: Transformers for Multimodal Self-Supervised Learning from Raw Video, Audio and Text.” *Proceedings of the Neural Information Processing Systems (NeurIPS)*, Virtual, December 2021.
- [C62] T.-Y. Pan, C. Zhang, Y. Li, H. Hu, D. Xuan, S. Changpinyo, **B. Gong**, and W.-L. Chao. “On Model Calibration for Long-Tailed Object Detection and Instance Segmentation.” *Proceedings of the Neural Information Processing Systems (NeurIPS)*, Virtual, December 2021.
- [C61] M. A. Jamal*, L. Wang, and **B. Gong**. “A Lazy Approach to Long-Horizon Gradient-Based Meta-Learning.” *Proceedings of the International Conference on Computer Vision (ICCV)*, Virtual, October 2021.
- [C60] C. Zhang, T.-Y. Pan, Y. Li, H. Hu, D. Xuan, S. Changpinyo, **B. Gong**, and W.-L. Chao. “MosaicOS: A Simple and Effective Use of Object-Centric Images for Long-Tailed Object Detection.” *Proceedings of the International Conference on Computer Vision (ICCV)*, Virtual, October 2021.
- [C59] X. Zhao, R. Vemulapalli, PA Mansfield, **B. Gong**, B. Green, L. Shapira, and Y. Wu. “Contrastive Learning for Label Efficient Semantic Segmentation.” *Proceedings of the International Conference on Computer Vision (ICCV)*, Virtual, October 2021.
- [C58] JW Shin, HB Lee, **B. Gong**, and SJ Hwang. “Large-Scale Meta-Learning with Continual Trajectory Shifting.” *Proceedings of the International Conference on Machine Learning (ICML)*, Virtual, July 2021.
- [C57] D. Kondratyuk, L. Yuan, Y. Li, L. Zhang, M. Tan, M. Brown, and **B. Gong**. “MoViNets: Mobile Video Networks for Efficient Video Recognition.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Virtual, June 2021.
- [C56] Y. Li*, X. Jia, R. Sang, Y. Zhu, B. Green, L. Wang, and **B. Gong**. “Ranking Neural Checkpoints.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Virtual, June 2021.
- [C55] X. Chen, C. Xie, M. Tan, L. Zhang, C.J. Hsieh, and **B. Gong**. “Robust and Accurate Object Detection via Adversarial Learning.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Virtual, June 2021.
- [C54] L. Yi, **B. Gong**, and T. Funkhouser. “Complete & Label: A Domain Adaptation Approach to Semantic Segmentation of LiDAR Point Clouds.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Virtual, June 2021.
- [C53] X. Fan, Q. Wang, J. Ke, F. Yang, **B. Gong**, and M. Zhou. “Adversarially Adaptive Normalization for Single Domain Generalization.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Virtual, June 2021.
- [C52] R. Qian, T. Meng, **B. Gong**, M.H. Yang, H. Wang, S. Belongie, and Y. Cui. “Spatiotemporal Contrastive Video Representation Learning.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Virtual, June 2021.
- [C51] Y. Ding, L. Wang, and **B. Gong**. “Analyzing Deep Neural Network’s Transferability via Frechet Distance.” *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, Remote, January 2021.

- [C50] Y. Li*, D. Huang, D. Qin, and **B. Gong**. “Improving Object Detection with Selective Self-Supervised Self-training.” *Proceedings of the European Conference on Computer Vision (ECCV)*, Remote, August 2020.
- [C49] M. Jamal*, M. Brown, L. Wang, M.H. Yang, and **B. Gong**. “Rethinking Class-Balanced Methods for Long-Tailed Visual Recognition from a Domain Adaptation Perspective.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Seattle, WA, June 2020.
- [C48] D. Wang-, Y. Li=*, L. Wang, and **B. Gong**. “Neural Networks Are More Productive Teachers Than Human Raters: Active Mixup for Data-Efficient Knowledge Distillation from a Blackbox Model.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Seattle, WA, June 2020.
- [C47] Z. Liu-, Z. Miao-, X. Pan, X. Zhan, D. Lin, S. Yu, and **B. Gong**. “Open Compound Domain Adaptation.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Seattle, WA, June 2020.
- [C46] C. Xie, M. Tan, **B. Gong**, J. Wang, A. Yuille, and Q.V. Le. “Adversarial Examples Improve Image Recognition.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Seattle, WA, June 2020.
- [C45] Y. Zhang=*, Z. Zhou-, P. David, X. Yue, Z. Xi, **B. Gong**, and H. Foroosh. “PolarNet: An Improved Grid Representation for Online LiDAR Point Clouds Semantic Segmentation.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Seattle, WA, June 2020.
- [C44] R. Zhai, C. Dan, D. He, H. Zhang, **B. Gong**, P. Ravikumar, C.-J. Hsieh, and L. Wang. “MACER: Attack-free and Scalable Robust Training via Maximizing Certified Radius.” *Proceedings of the International Conference on Learning Representations (ICLR)*, Addis Ababa Ethiopia, April 2020.
- [C43] C. Gan-, Y. Zhang-, J. Wu, **B. Gong**, and J. Tenenbaum. “Look, Listen, and Act: Towards Audio-Visual Embodied Navigation.” *Proceedings of the International Conference on Robotics and Automation (ICRA)*, Paris, France, June 2020.
- [C42] Z. Yang, **B. Gong**, L. Wang, W. Huang, D. Yu, and J. Luo. “A Fast and Accurate One-Stage Approach to Visual Grounding.” *Proceedings of the International Conference on Computer Vision (ICCV)*, Seoul, Korea, October 2019.
- [C41] X. Yue, Y. Zhang, S. Zhao, A. Sangiovanni-Vincentelli, K. Keutzer, and **B. Gong**. “Domain Randomization and Pyramid Consistency: Simulation-to-Real Generalization without Accessing Target Domain Data.” *Proceedings of the International Conference on Computer Vision (ICCV)*, Seoul, Korea, October 2019.
- [C40] Q. Lian, F. Lv, L. Duan, and **B. Gong**. “Constructing Self-motivated Pyramid Curriculums for Cross-Domain Semantic Segmentation: A Non-Adversarial Approach.” *Proceedings of the International Conference on Computer Vision (ICCV)*, Seoul, Korea, October 2019.
- [C39] G. Shen, W. Huang, C. Gan, M. Tan, J. Huang, W. Zhu, and **B. Gong**. “Facial Image-to-Video Translation by a Hidden Affine Transformation.” *Proceedings of the 27th ACM international conference on Multimedia (MM)*, Nice, France, October 2019.
- [C38] Y. Li*, L. Li, L. Wang, T. Zhang, and **B. Gong**. “NATTACK: Learning the Distributions of Adversarial Examples for an Improved Black-Box Attack.” *Proceedings of the International Conference on Machine Learning (ICML)*, Long Beach, CA, June 2019.
- [C37] Z. Liu, Z. Miao, X. Zhan, J. Wang, **B. Gong**, and S. Yu. “Large-scale Long-Tailed Recognition in an Open World.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Long Beach, CA, June 2019.
- [C36] J. Shi, J. Xu, **B. Gong**, and C. Xu. “Not All Frames Are Equal: Weakly-Supervised Video Grounding with Contextual Similarity and Visual Clustering Losses.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Long Beach, CA, June 2019.

- [C35] X. Tang, **B. Gong**, Y. Yu, H. Yao, Y. Li, H. Xie, and X. Wang. “Joint Modeling of Dense and Incomplete Trajectories for Citywide Traffic Volume Inference.” *Proceedings of The Web Conference (WWW)*, San Francisco, CA, May 2019.
- [C34] Y. Zhang*, H. Foroosh, P. David, and **B. Gong**. “CAMOU: Learning Physical Vehicle Camouflages to Adversarially Attack Detectors in the Wild.” *Proceedings of The International Conference on Learning Representations (ICLR)*, New Orleans, LA, May 2019.
- [C33] M. Fang, C. Zhou, B. Shi, **B. Gong**, J. Xu, and T. Zhang. “DHER: Hindsight Experience Replay for Dynamic Goals.” *Proceedings of The International Conference on Learning Representations (ICLR)*, New Orleans, LA, May 2019.
- [C32] D. Zhu, Z. Lin, X. Wang, **B. Gong**, and T. Yang. “A Robust Zero-Sum Game Framework for Pool-based Active Learning.” *Proceedings of The International Conference on Artificial Intelligence and Statistics (AISTATS)*, Naha, Japan, April 2019.
- [C31] L. Fan, W. Huang, C. Gan, J. Huang, and **B. Gong**. “Controllable Image-to-Video Translation: A Case Study on Facial Expression Generation.” *Proceedings of The Thirty-Third AAAI Conference on Artificial Intelligence (AAAI)*, Hawaii, January 2019.
- [C30] L. Li* and **B. Gong**. “End-to-End Video Captioning with Multitask Reinforcement Learning.” *Proceedings of the IEEE Winter Conference on Applications of Computer Vision (WACV)*, Hawaii, January 2019.
- [C29] Z. He, **B. Gong**, and D. Fan. “Optimize Deep Convolutional Neural Network with Ternarized Weights and High Accuracy.” *Proceedings of the IEEE Winter Conference on Applications of Computer Vision (WACV)*, Hawaii, January 2019.
- [C28] H. Hu=, L. Chen=, **B. Gong**, and F. Sha. “Synthesize Policies for Transfer and Adaptation across Environments and Tasks.” *Proceedings of the Neural Information Processing Systems (NeurIPS)*, Montreal, Canada, December 2018.
- [C27] Y. Li*, L. Wang, T. Yang, and **B. Gong**. “How Local is the Local Diversity? Reinforcing Sequential Determinantal Point Processes with Dynamic Ground Sets for Supervised Video Summarization.” *Proceedings of the European Conference on Computer Vision (ECCV)*, Munich, Germany, September 2018.
- [C26] A. Sharghi*, A. Borji, C. Li, T. Yang, and **B. Gong**. “Improving Sequential Determinantal Point Processes for Supervised Video Summarization.” *Proceedings of the European Conference on Computer Vision (ECCV)*, Munich, Germany, September 2018.
- [C25] MA. Jamal*, H. Li, and **B. Gong**. “Face Detector Adaptation without Negative Transfer or Catastrophic Forgetting.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Salt Lake City, Utah, June 2018.
- [C24] L. Fan*=, W. Huang=, C. Gan, S. Ermon, **B. Gong**, and J. Huang. “End-to-End Learning of Motion Representation for Video Understanding.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Salt Lake City, Utah, June 2018.
- [C23] C. Gan*, **B. Gong**, H. Su, and L. Guibas. “Geometry-Guided CNN for Self-Supervised Video Representation Learning.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Salt Lake City, Utah, June 2018.
- [C22] X. Wei*=, **B. Gong**=, Z. Liu, W. Lu, and L. Wang. “Improving the Improved Training of Wasserstein GANs: A Consistency Term and Its Dual Effect.” *Proceedings of the International Conference on Learning Representations (ICLR)*, Vancouver Canada, April 2018.
- [C21] Y. Ding*, L. Wang, D. Fan, and **B. Gong**. “A Semi-Supervised Two-Stage Approach to Learning from Noisy Labels.” *Proceedings of the IEEE Winter Conference on Applications of Computer Vision (WACV)*, Lake Tahoe, NV, March 2018.
- [C20] Z. Yang*, **B. Gong**, and S. Narayanan. “Weighted Geodesic Flow Kernel for Interpersonal Mutual Influence Modeling and Emotion Recognition in Dyadic Interactions.” *Proceedings of*

the International Conference on Affective Computing and Intelligent Interaction (**ACII**), San Antonio, TX, October 2017.

- [C19] Y. Zhang*, P. David, and **B. Gong**. “Curriculum Domain Adaptation for Semantic Segmentation of Urban Scenes.” *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, Venice, Italy, October 2017.
- [C18] C. Gan*, Y. Li*, H. Li, C. Sun, and **B. Gong**. “VQS: Linking Segmentations to Questions and Answers for Supervised Attention in VQA and Question-Focused Semantic Segmentation.” *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, Venice, Italy, October 2017.
- [C17] A. Sharghi*, J. Laurel*, and **B. Gong**. “Query-Focused Video Summarization: Dataset, Evaluation, and A Memory Network Based Approach.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Honolulu, Hawaii, Jun. 2017.
- [C16] M. Kalayeh*, **B. Gong**, and M. Shah. “Improving Facial Attribute Prediction using Semantic Segmentation.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Honolulu, Hawaii, Jun. 2017.
- [C15] Z. Li*, **B. Gong**, and T. Yang. “Improved Dropout for Shallow and Deep Learning.” *Proceedings of the Neural Information Processing Systems (NIPS)*, Barcelona, Spain, Dec. 2016.
- [C14] C. Gan*, C. Sun, L. Duan, and **B. Gong**. “Labeling-Free Video Recognition by Mutually Voting for Relevant Web Images and Web Video Frames.” *Proceedings of the European Conference on Computer Vision (ECCV)*, Amsterdam, Netherlands, Oct. 2016.
- [C13] A. Sharghi*, **B. Gong**, and M. Shah. “Query-Focused Extractive Video Summarization.” *Proceedings of the European Conference on Computer Vision (ECCV)*, Amsterdam, Netherlands, Oct. 2016.
- [C12] W-L. Chao-, S. Changpinyo-, **B. Gong**, and F. Sha. “An Empirical Study and Analysis of Generalized Zero-Shot Learning for Object Recognition in the Wild.” *Proceedings of the European Conference on Computer Vision (ECCV)*, Amsterdam, Netherlands, Oct. 2016.
- [C11] Y. Zhang*, **B. Gong**, and M. Shah. “Fast Zero-Shot Image Tagging.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Las Vegas, NV, Jun. 2016.
- [C10] C. Gan*, T. Yang, and **B. Gong**. “Learning Attributes Equals Multi-Source Domain Generalization.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Las Vegas, NV, Jun. 2016.
- [C9] S. Changpinyo-, W. Chao-, **B. Gong**, and F. Sha. “Synthesized Classifiers for Zero-Shot Learning.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Las Vegas, NV, Jun. 2016.
- [C8] W. Chao-, **B. Gong**-, F. Sha, and K. Grauman. “Large-Margin Determinantal Point Processes.” *Proceedings of the Conference on Uncertainty in Artificial Intelligence (UAI)*, Amsterdam, Netherlands, July 2015. [\[Link\]](#)
- [C7] **B. Gong**-, W. Chao-, K. Grauman, and F. Sha. “Diverse Sequential Subset Selection for Supervised Video Summarization.” *Proceedings of the Neural Information Processing Systems (NIPS)*, Montreal, Canada, Dec. 2014.
- [C6] **B. Gong**, K. Grauman, and F. Sha. “Reshaping Visual Datasets for Domain Adaptation.” *Proceedings of the Neural Information Processing Systems (NIPS)*, Lake Tahoe, NV, Dec. 2013.
- [C5] **B. Gong**, K. Grauman, and F. Sha. “Connecting the Dots with Landmarks: Discriminatively Learning Domain-Invariant Features for Unsupervised Domain Adaptation.” *Proceedings of the International Conference on Machine Learning (ICML)*, Atlanta, GA, Jun. 2013.
- [C4] **B. Gong**, Y. Shi, F. Sha, and K. Grauman. “Geodesic Flow Kernel for Unsupervised Domain Adaptation.” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Providence, RI, Jun. 2012.

- [C3] **B. Gong**, J. Liu, X. Wang, and X. Tang. “3D Object Retrieval with Semantic Attributes.” *Proceedings of the 19th ACM international conference on Multimedia (ACM MM)*, Scottsdale, Arizona, Dec. 2011. (demo)
- [C2] **B. Gong**, C. Xu, J. Liu, and X. Tang. “Boosting 3D Object Retrieval by Object Flexibility”. *Proceedings of the 17th ACM international conference on Multimedia (ACM MM)*, Beijing, China, Oct. 2009.
- [C1] **B. Gong**, Y. Wang, J. Liu, and X. Tang. “Automatic Facial Expression Recognition on a Single 3D Face by Exploring Shape Deformation”. *Proceedings of the 17th ACM international conference on Multimedia (ACM MM)*, Beijing, China, Oct. 2009.

SELECTED PEER-REVIEWED WORKSHOP PUBLICATIONS

- [W3] D. Kondratyuk, M. Tan, M. Brown, and **B. Gong**. “When Ensembling Smaller Models is More Efficient than Single Large Models.” *The 4th Workshop on Visual Understanding by Learning from Web Data* at CVPR, June 2020.
- [W2] A. Mazaheri*, B. Gong, and M. Shah. “Learning a Multi-Concept Video Retrieval Model with Multiple Latent Variables.” *The 12th IEEE International Workshop on Multimedia Information Processing and Retrieval*, Dec. 2016.
- [W1] **B. Gong**, F. Sha, and K. Grauman. “Overcoming Dataset Bias: An Unsupervised Domain Adaptation Approach.” *The First International Workshop on Large Scale Visual Recognition and Retrieval (BigVision)* at NIPS, Lake Tahoe, NV, Dec. 2012.

PATENTS

US20250245800A1: Robust automatic image and video quality assessment, 2025.

US12395685B2: Highly efficient model for video quality assessment, 2025.

US20240320493A1: Improved two-stage machine learning for imbalanced datasets, 2024.

US20240119307A1: Personalized federated learning via sharable basis models, 2023.

US20240096076A1: Semantic segmentation neural network for point clouds, 2022.

US10885345B2: End-to-end video captioning with multi-task reinforcement learning, 2019.

TEACHING EXPERIENCES

CS 599: Deep Visual Generative Models (BU)
Spring 2026, 18 students, Rating: **4.89/5** (Department mean: 4.21)

CS 542: Principles of Machine Learning (BU)
Fall 2025, 99 students, Rating: **4.28/5** (Department mean: 4.21)
Spring 2025, 39 students, Rating: **4.63/5**
Fall 2024, 91 students, Rating: **4.43/5**

DATA 2040: Deep Learning and Special Topics in Data Science at Brown University
Spring 2021, 40 students, Rating: **4.33/5** (Department median: 3.94)

CAP 4453: Robot Vision (UCF)
Fall 2016, 64 students, Rating: 3.90/5 (Department median: 3.81)
Fall 2017 (fully online), 30 students, Rating: **4.29/5** (Department median: 3.82)

CAP 6412: Advanced Computer Vision (UCF)
Spring 2016, 18 students, rating: **4.10/5** (Department median: 3.85)

STUDENTS

Ph.D. students at BU:

Max Whitton (Dean’s Fellow)	09/2025 –
Lingxiao Li	09/2025 –

Yuwen Tan (Hariri Institute Graduate Student Fellow)	09/2024 –
Wenqi Wang	09/2024 –
Zecheng Wang (Dean’s Fellow)	09/2024 –
Shengao Wang (co-supervised with Venkatesh Saligrama)	09/2024 –

Ph.D. graduates at UCF:

Yang Zhang	08/2015 – 12/2017; co-supervised with Hassan Foroosh, 2018 – 2020
Aidean Sharghi	08/2015 – 12/2017
Abdullah Jamal	01/2016 – 12/2017; co-supervised with Liqiang Wang, 2018 – 2021
Yifan Ding	01/2016 – 12/2017; co-supervised with Liqiang Wang, 2018 – 2023
Yandong Li	08/2017 – 12/2017; co-supervised with Liqiang Wang, 2018 – 2021
Samer Iskander	co-supervise with Dr. Niels Lobo, 01/2016 – 05/2016

Visiting students:

BU: Logan Lawrence (from UMass Amherst, 2026), Yuan Qing (from Columbia University, 2025–26), Fuhang Kuang (from Tsinghua, 2025)
UCF: Lijun Li (from Beihang University, 2016–17)

Master students:

BU: Harsha Beth (2026), Mohamed Ifreen Seyed Ibrahim (2025)
UCF: Fareeha Irfan (2015–17), Suhas Nithyanand (2016), Rohan Singh Rajput (2016)

Undergraduate students:

Michael Wakeham (BU)	2025 – 2026
Arjun Chandra (BU)	2025 – 2026
Joey Huang (BU)	2024 – 2025
Adam Vest, Univ. of Louisville (NSF Research Experiences for Undergraduates (REU))	2017
Geraldine Versfeld, University of Central Florida (NSF REU)	2017
Truman Thames, Fayetteville State University UNC (NSF REU)	2017
Jacob Scott Laurel, University of Alabama at Birmingham (NSF REU)	2016
Kylie McCarty, University of Central Florida (NSF REU)	2016
Kevin Duarte, University of Central Florida (NSF REU)	2016
Michael Lopez (undergraduate research program)	Spring 2016
Adam Albright, University of Central Florida (senior design)	2016 – 2017
Qiang Li, University of Central Florida (senior design)	2016 – 2017
Kyle Ferguson, University of Central Florida (senior design)	2016 – 2017

Ph.D. defense and candidacy committee member for

Expected to graduate in 2027: Yuanhao Ban (University of California at Los Angeles), Aoming Liu (BU), Sherin Muckatira (UMass Lowell)
Graduated in 2026: Fangrui Zhu (Northeastern University), Piotr Teterwak (BU), Arijit Ray (BU), Chau Pham (BU), Sina Malakouti (University of Pittsburgh), Jeffrey Willette (Korea Advanced Institute of Science and Technology)
Graduated in 2025: Kaihong Wang (BU), Zheyuan Zhang (Northwestern University), Zihang Zou (UCF)
Graduated in 2024: Bruno Andreis (Korea Advanced Institute of Science and Technology), Minseon Kim (Korea Advanced Institute of Science and Technology)
Graduated in 2023: Zixiang Zhou (UCF), Yifan Ding (UCF), Dongdong Wang (UCF)
Graduated in 2022: Hae Beom Lee (Korea Advanced Institute of Science and Technology)
Graduated in 2021: Muhammad Abdullah Jamal (UCF), Yandong Li (UCF)

Graduated in 2020: Yang Zhang (UCF)
 Graduated in 2018: Maryam Jaber (UCF), Dustin Morley (UCF)
 Graduated in 2017: Uzair Tariq (UCF), Hong Zhang (UCF)
 Graduated in 2016: Kenneth Thompson (UCF)

Master and Honors Undergraduate defense and candidacy committee member for	
Chaitanya Chakka, Boston University	Graduated in 2026
Sujayanand Kingsly, Boston University	Graduated in 2026
Yitong Wang, Boston University	Graduated in 2025
Seunghwan Hyun, Boston University	Graduated in 2025

GRANTS
& AWARDS

NSF CAREER: *Democratizing the Pretraining of Vision Foundation Models: A Developmentally Plausible Framework*

\$599,999, Sole Principle Investigator (PI), IIS-2540851	09/2026 – 08/2031
--	-------------------

Awards with research funding (\$140,000 in total)

\$100,000, Sony Faculty Innovation Award	09/2025 – 08/2026
\$30,000, Faculty of Computing & Data Science Junior Faculty Fellow	08/2025 – 08/2028
\$10,000, Rafik B. Hariri Institute Junior Faculty Fellow	05/2026 – 05/2029

Unrestricted gifts (\$202,000 in total)

\$30,000 from Creatify Inc.	11/2025
\$10,000 from Adobe Inc.	05/2025
\$15,000 from SNAP Inc.	04/2025
\$30,000 from AirBrush Innovation Inc.	04/2025
\$100,000 from FutureWei Technologies Inc. (UCF, Declined)	07/2017
\$10,000 from Adobe Research (UCF)	05/2017
\$7,000 from Adobe Research (UCF)	10/2016

Cloud credits (\$90,000 in total)

\$50,000 from Google Gemini Credits	12/2025
\$15,000 from Microsoft Research	04/2025
\$25,000 from Rafik B. Hariri Institute AI Research Resource Program	12/2024

NSF BIGDATA: IA: *Distributed Semi-Supervised Training of Deep Models and Its Applications in Video Understanding*

\$662,431 + \$42,500 cloud credits, PI, IIS-1741431	09/2017 – 08/2020
The first of its kind ever granted to UCF. Transferred to ex-colleagues after leaving UCF	

NSF CRII: RI: *Multi-Source Domain Generalization Approaches to Visual Attribute Detection*

\$175,000, Sole PI, IIS-1566511	05/2016 – 04/2018
The first of its kind ever granted to UCF	